RCI Agent demonstration manual

Evidence desk presenter guide | 28 October 2026

Editions: English PDF, Korean Markdown, Korean PDF, Japanese Markdown, Japanese PDF, Chinese Markdown, Chinese PDF.

1. Demo at a glance

Audience: presenters, analysts, and stakeholders exploring the implemented investigation demo. Allow 15-20 minutes after setup. This manual describes the current app, not every capability in the aspirational scenario pack.

Opening message: "This is a fictional investigation workbench. The agent helps organize and challenge evidence; the analyst controls amendments and review. No bank, case-management system, or regulatory decision service is connected."

What the audience will see

Segment	Time	Demonstration
Open and prepare	3 min	Open an alert from the alert list, confirm the question, then collect packaged fictional evidence.
Inspect and challenge	4 min	Inspect source limitations and ask the configured agent to challenge findings.
Amend and clarify	4 min	Save an analyst correction and submit a bounded scope form.
Review and prepare	4 min	Review six outputs and prepare a copyable case note.
Change and revisit	3 min	Invalidate the review, inspect history, and reopen saved work.

The case in one minute

Harbour Components Ltd is a fictional industrial component wholesaler. Existing alert **ALT-2041**, shown in **CASE-1041**, has original AML priority **High**. Three receipts total **GBP 48,000**, followed by **GBP 45,500** to Meridian Supply: **94.79%** of the receipts, **56 minutes** after the first incoming payment.

The beneficiary is not observed in the available 90-day history. That is a stipulated scenario fact, not a deduction from the displayed rows or proof of a new commercial relationship. A supplier-payment explanation remains plausible. The invoice is missing and a previous case is restricted.

Choose the presentation route

Full route: follow sections 2-8. Preparation uses no model inference, but sending messages and submitting the scope form call the configured model and may incur charges. Model wording, timing, and proposed updates are not guaranteed.

No-inference route: use preparation, source inspection, manual amendments, review, history, and reopening. Skip sending agent prompts, policy requests, and scope-form submission. Demonstrate the form by opening and cancelling it. Explicitly say that this route does not demonstrate live agent analysis.

Core distinction: saved amendments become context for subsequent agent turns; they do not automatically rewrite every other stage's prose or modify original source records. Affected stages are marked for re-review.

2. Prepare the environment

Before the audience arrives

Use the repository root in Windows PowerShell. Follow **README.md** for first-time installation, dependency restoration, model configuration, and authentication. Do not copy credentials into prompts or demonstration material.

If upgrading, stop the app's API and worker and back up the configured data directory, `.data` by default, before startup migrations. Restore the backend and frontend locks as documented in the README. Restart both API and worker after extension or dependency changes; their configuration and data directory must match.

Select a model with authorized inference access for the full route. Select it before beginning. Opening an alert creates its one linked case and reopening it resumes that case, so rehearsal does not create duplicates. Each **New investigation** creates another unlinked CASE-1041 session with the same title, so use it sparingly.

Start the three processes

Terminal 1 - API

```
.\.venv\Scripts\python.exe -m uvicorn skill_chat.api.app:create_app --factory --host 127.0.0.1 --port 8000
```

Terminal 2 - durable worker

```
.\.venv\Scripts\python.exe -m skill_chat.tasks.worker
```

Terminal 3 - frontend

```
npm --prefix .\web run dev -- --host 127.0.0.1 --port 5173
```

Open **http://127.0.0.1:5173**. Use the exact configured frontend origin. If a port is already occupied, use the README's alternate-port instructions rather than stopping an unknown process.

Check readiness

```
Invoke-RestMethod http://127.0.0.1:8000/health/live
Invoke-RestMethod http://127.0.0.1:8000/health/ready
```

Readiness requires a recent worker heartbeat. It does not validate cloud permissions or model quality. Rehearse one short model request separately before a full-route presentation; it may be billable.

Presenter checklist: use fictional data only; enlarge the browser for the audience; locate the model picker; keep the API and worker terminals available for diagnosis without projecting logs or secrets. Use the ordinary app URL above, not the in-memory layout comparison at `?prototype=centered-chat`. Do not use the isolated automated-test servers as a substitute for the normal app.

3. Open the desk and inspect evidence

Open the alert

1. In the sidebar **Alerts (fictional demo)** list, review the four alerts. Each row shows the alert ID, priority, subject, trigger, and either **No case yet. Review by** a date or **Case open at** a stage. The review-by date is a reminder, not regulatory advice.
2. Choose **Open alert ALT-2041, Harbour Components Ltd**. This creates the one case linked to that alert. Choosing the same alert later (its button then reads **Resume alert**) reopens that case instead of creating a duplicate.
3. Confirm **CASE-1041 evidence desk**, the **Demo connector / fictional records** label, and the original **High** alert priority. Nothing has been prepared yet: the desk shows the starting **Question to answer** and states that no evidence has been collected.
4. Point out the six stages: Detect, Investigate, Gather evidence, Analyze, Human review, and Decision. The case opens at **Investigate**. Click **Detect** to show that it is **Read-only**: it records the original alert and priority, starts acknowledged, is not invalidated by later changes, and cannot be amended. **Continue to Investigate** returns to the working stage.
5. Optionally use **Clarify investigation scope** before preparing evidence (section 5).
6. Click **Prepare investigation**. Wait for the preparation task to succeed and **Refresh demo evidence** to become enabled.

Expected result: stage outputs update. Select the central **Evidence** tab to see the populated transaction timeline and evidence register. Preparation retrieves packaged fixture records through the backend worker without model inference. It is not a live financial-system query.

Alert without evidence (optional, 2 min): open another alert, for example **ALT-2038**, and click **Prepare investigation**. Only ALT-2041 has packaged demonstration evidence, so preparation records one missing source (**SRC-00**) and the gap **GAP-01**. Analyze states that no finding can be drawn. Use this to show that the desk reports

a gap rather than inventing evidence.

Referral path: New investigation remains below the alert list for a referral or manual lead. It creates an independent CASE-1041 session that is not linked to any alert, and its Detect stage stays editable.

Explain navigation

The implemented **A layout** has **Investigation sessions** on the left, the workflow and stage workspace in the center, and **Investigation assistant** with conversation and composer on the right.

Area or control	Presenter action
Output tab	Edit the inspected stage, amend shared facts, and open revision history.
Evidence tab	Inspect the transaction timeline, evidence register, and source records.
Review tab	Enter analyst rationale, review the note, and prepare or copy it. **Review & record** opens this tab too.
Sidebar dividers	On desktop, drag either divider. Keyboard users can focus it and use Left/Right arrows; Home/End select its minimum/maximum width.
Chat	Open or close the right conversation sidebar. Its close icon returns to the workspace without discarding the message draft.
Sun/moon button	Switch light/dark mode. The first visit follows the system; an explicit choice is saved when browser storage is available.

On screens up to 1,000 pixels wide, history and chat are drawers rather than simultaneous columns; resizing handles are hidden. Use **Toggle history** to open sessions and **Chat** to open the assistant. Close the chat drawer before working in the central area.

Click **Analyze** to inspect its **Output** tab. Clicking any stage returns to Output; it does not advance the workflow. Use **Make current stage** only when deliberately changing the current workflow stage. Reviewing a note later sets Human review; preparing a note sets Decision.

Inspect the underlying records

Select **Evidence**, then open **Inspect SRC-01** from a relevant evidence item. Read the selected ledger records, fixture snapshot, version, retrieval timestamp, and coverage. Use **Close source** before moving on.

Record	UTC	Amount	Counterparty
TX-781	09:12	+GBP 18,000	Alder Industrial
TX-782	09:31	+GBP 16,000	Beacon Assembly
TX-783	09:54	+GBP 14,000	Cedar Parts
TX-784	10:08	-GBP 45,500	Meridian Supply

Select **Analyze** to return to its **Output** tab and point to the 94.79% comparison and 56-minute window. Explain that selected records do not supply balances or establish fund lineage or intent.

Return to **Evidence**. Inspect **SRC-04** to show the restricted prior case and **SRC-07** to show the missing invoice. **All sources (7)** provides another way to open each source.

Expected result: SRC-04 remains restricted and SRC-07 remains missing. Neither means an adverse finding or a completed corroboration check. Refreshing evidence changes retrieval observations; it does not make unavailable records accessible or make old fixture content newly current.

Talk track: "Source freshness and retrieval time are different. We can see what was supplied, what was not supplied, and why the conclusion remains limited."

4. Challenge and amend an interpretation

Ask the configured agent

For the full route, select **Analyze** in the center. In the right **Investigation assistant** sidebar, click **Draft a challenge**. If the assistant is closed, open it with **Chat** first. This only populates the composer. Read the prompt, then click **Send message**.

Wait for this new task to finish; do not mistake a preceding successful task for the new result. Inspect the response and stage output rather than promising a fixed answer.

Suggested follow-up prompt:

> What evidence challenges the concern? Separate observed facts, calculated comparisons, alternative explanations, and gaps. Propose a revised Analyze output with evidence references. Do not infer wrongdoing, fund lineage, or the contents of the restricted case.

Expected behavior: the agent receives the current saved investigation context and can propose scoped output updates. A proposal becomes canonical only after successful task completion. A response is not proof that an output was changed: inspect the stage's text and version. If no update was proposed, ask explicitly or use the manual amendment below.

Save an analyst correction

Select **Analyze**, replace **Stage output** with the following example, and enter the reason.

> Selected records show GBP 48,000 incoming and GBP 45,500 outgoing within 56 minutes [E-01]. Timing alone does not establish intent or fund lineage. Supplier payments fit the business profile [E-03], but the invoice remains unverified [E-04, GAP-01]. Restricted case contents are not adverse evidence [GAP-02].

Reason for stage amendment:

> Preserve the alternative explanation and distinguish observed facts from interpretation.

Before saving, click **Detect** and return to **Analyze**. Then switch to **Evidence** and back to **Output**. The unsaved output and reason should remain through both kinds of navigation. Click **Apply stage amendment** to persist them.

Expected result: the revision and output version increase; Analyze and downstream outputs require review; a prepared note is cleared. Earlier stage prose and the original source records are not overwritten.

Explain persistence precisely

Saved edits enter subsequent agent context. Unsaved stage edits, reasons, and review rationale survive tab changes within the mounted desk; stage edits also survive stage inspection. A message draft survives closing and reopening the assistant. None of these unsaved drafts is a saved record; do not rely on them across reloads or session changes.

In **Output**, open **Revision history** to inspect the amendment, reason, **Previous text**, and **Saved text** where available. Use **Load older revisions** when offered. This is application revision history, not proof of an external compliance audit system.

5. Clarify scope and check fictional policy

Demonstrate the bounded A2UI form

1. With no active task, click **Clarify investigation scope** below the active central tab's content; scroll down if needed. On mobile, close the chat drawer first.
2. In **Investigation scope**, enter the example below.
3. For the full route, click **Apply scope and ask agent**.
4. Wait for **Scope clarification resolved.** and for the newly created task to succeed before continuing.

> Assess the permitted supplier-payment evidence. Keep the missing invoice and restricted prior case explicit. Do not infer a new commercial relationship from absence in the available 90-day history.

Expected behavior: the answer is persisted and queues one ordinary durable agent continuation. The worker does not wait while the analyst types. Pending, resolved, and cancelled status survives reopening.

For the no-inference route, do not submit. Click **Cancel scope clarification** instead. Opening or cancelling the fixed form is not an agent analysis demonstration.

If the official renderer cannot display the form, an accessible **Plain form fallback** may appear. It supports the same bounded answer and submission; it is not permission to load arbitrary agent-generated HTML or screens.

Optional policy request

In **Evidence**, open **All sources (7)** and inspect **SRC-05** and **SRC-06**. These contain fictional **G-AML-07 v3** and **L-01 v2** passages.

In the right assistant sidebar, click **Draft a policy check**, inspect the prepared composer text, and send it only on the full route. Alternatively use:

> Using the supplied fictional G-AML-07 v3 and L-01 v2 passages, identify the checks and gaps that remain before a human could recommend closure or escalation. Cite the passage locators. Distinguish this demo policy from law and do not claim the invoice has been corroborated.

Expected discussion: supporting and contradicting evidence must be considered; investigator validation and rationale remain necessary; the missing invoice and inaccessible case remain limitations. The generated wording is model-dependent and must be checked against the source passages.

Optional AML skills

The six workflow stages do not automatically run six skills. If demonstrating an installed AML skill on the full route, prepare the evidence first, then choose one from **Skills** in the assistant composer or type its slash command. The popup shows names only.

Skill	Demonstration focus
`/aml-case-context`	Explain the recorded alert, customer context, and information gaps.
`/aml-transaction-behaviour`	Assess verified transaction indicators against expectations and history.
`/aml-hypothesis-review`	Test concerns against legitimate explanations and counterevidence.
`/aml-policy-applicability`	Identify applicable supplied procedures and human approval requirements.
`/aml-investigation-readiness`	Identify incomplete checks; missing required evidence must block a ready result.
`/aml-case-narrative`	Draft a cited, editable case note for analyst review.

These are optional instructions, not a validated risk model, approved real bank policy pack, calculation engine, or live case integration. Calculations must come from verified preparation metrics or identified approved tools, not language-model arithmetic. Selecting a skill alone does not run inference; sending its prompt does. A selected skill remains active for follow-ups; use **Clear skill** on its composer chip to return to ordinary prompts.

Explain AG-UI without overclaiming

AG-UI delivers replayable investigation snapshots and clarification activity to the fixed React workspace. User actions are validated server-side commands; conversation text keeps its existing separate stream.

A2UI renders only the restricted scope form. This is not arbitrary agent-authored workflow design, a complete AG-UI text/tool lifecycle implementation, or SDK checkpoint/resume of a worker waiting for human input.

6. Review and prepare the case note

Complete agent turns and amendments before this segment. Each further material change may invalidate review.

Fill the human-owned outputs

Preparation populates evidence and analysis, not a complete human judgment. Inspect every stage. Do not acknowledge generic "Awaiting investigation preparation and human assessment" text as a completed assessment.

For **Human review**, amend the output as appropriate to the demonstration:

> Analyst checked the selected transaction comparison and considered the plausible supplier explanation. The invoice is unavailable [GAP-01]; prior case contents remain restricted [GAP-02]. No conclusion about intent or fund lineage is established.

For **Decision**, record a proposed disposition, not an executed decision:

> Proposed next action: obtain invoice corroboration and request authorized access to the prior case before considering closure. No closure, escalation approval, regulatory filing, or account restriction has been executed.

Enter a meaningful **Reason for stage amendment** for each change and click **Apply stage amendment**. These examples are presenter-supplied judgments, not automatic agent findings.

Acknowledge and review

1. Visit the five editable stage outputs after the final amendment. Detect is read-only and already acknowledged for alert-opened cases.
2. Validate each output and its evidence references. Click **Acknowledge this output** only when satisfied.
3. Wait until each editable stage shows **Acknowledged**, not merely a disabled button while a request is running.
4. Select the central **Review** tab, or click **Review & record**. In **Analyst rationale**, enter the example below, adjusting it to the actual outputs.
5. Click **Mark note reviewed**, then **Prepare case note**.

> I checked the selected records and alternative supplier explanation. The invoice gap and restricted prior case remain unresolved. Further authorized evidence collection is proposed; no external decision or approval is confirmed.

Expected result: preparation must already have succeeded; all affected outputs must be acknowledged; rationale must be nonempty. The prepared note appears with the shared finding, stage outputs, rationale, evidence-derived references and gaps, and **External saving / approval: unconfirmed**.

Copy, do not claim external completion

Click **Copy case note**, or manually select and copy the displayed note if clipboard access is unavailable. Use a disposable demonstration document as the destination, not a real case system.

Talk track: "This prepares a record for the existing human workflow. Copying is not saving to another system, approving a disposition, or closing a case."

7. Show that changes invalidate review

Amend a shared fact

After preparing the note, select **Output** and open **Shared fact across all stages**. Change **Shared beneficiary description** to:

> Meridian Supply was not observed in the available 90-day history; the age of the commercial relationship is unknown.

In **Reason for shared fact amendment**, enter:

> Avoid equating absence from available history with a new commercial relationship.

Click **Apply shared fact amendment**.

Expected result: the shared finding changes; the five editable outputs require re-review (Detect stays acknowledged on alert-opened cases); the reviewed flag and prepared note are cleared. Return to **Review**: **Copy case note** is no longer available and note review is blocked until the required acknowledgments are renewed. Switching to Output alone only hides the note; the amendment is what invalidates it.

The shared wording is now available to all subsequent agent turns. It does not automatically replace each stage's text or alter immutable source content. If demonstrating agent propagation, send a short request to explain the amended shared fact and propose any needed output corrections; this is a model call.

Inspect the record and reopen

Return to **Output** and open **Revision history**. Show the actor, action, revision, reason, and before/after text for an amendment. Load older revisions when available.

Save any remaining draft before reloading. Reload the page and reopen the case by choosing **Resume alert ALT-2041** in the alert list, or select it from **Investigation sessions**. If several sessions share that title, select the one just used. Confirm the saved amendment, evidence, and clarification status remain.

Reloading does not itself rerun evidence preparation or agent analysis. Browser disconnection does not cancel durable background work.

Keep ordinary conversations distinct

Choose **New conversation** to show that ordinary chat, PDF review, optional skills, and saved HTML artifacts remain available. Do not introduce a removed historical RFP skill as a current product feature.

Close the presentation: "We have a persistent, editable evidence desk with traceable revisions. The analyst can challenge and correct the agent. Source limitations stay visible, and the app does not execute consequential financial or regulatory decisions."

Reset for the next demonstration

Open a different alert (for example ALT-2036) or use **New investigation** for a fresh fictional session; reopening ALT-2041 resumes the existing case, it does not reset it. To repeat ALT-2041 from the start, delete its session from **Investigation sessions** (this frees the alert for a new case). Do not delete `.data` to reset the demo. Avoid **Refresh demo evidence** as a reset button: it creates a new preparation turn and can invalidate prior review.

8. Troubleshooting and boundaries

Symptom	Presenter action
Task stays queued	Check worker readiness and matching API/worker configuration. Do not submit duplicates.
Controls are disabled	Wait for the active turn or mutation to finish. Review also needs prepared evidence, acknowledgments, and rationale.
Revision conflict	Keep the local draft. Compare with the saved output; use **Keep draft and use latest revision** when offered, then deliberately reapply with a reason. Do not blindly overwrite another revision.
Model fails or is unauthorized	Keep the failure visible. Check provider setup outside the presentation or switch to the no-inference route. Do not call a fixture response live inference.
Context exceeds budget	Shorten excessive saved outputs or choose a suitably configured model budget. A retry does not silently rebase stale analysis.
Source remains unavailable	Explain SRC-04 is restricted and SRC-07 is missing by design. Do not imply access bypass or successful invoice retrieval.
Note cannot be prepared	Finish preparation; amend placeholders; acknowledge every affected stage; enter rationale; mark the note reviewed.
Scope renderer fails	Use the plain fallback if offered. If the surface is invalid, reopen the session to reload it.
Copy fails	Select and copy the displayed note manually; external saving remains unconfirmed.
Sessions or assistant are hidden	Use **Toggle history** for sessions or **Chat** for the assistant. On narrow screens, close the chat drawer before editing the workspace.
Timeline, review, or history seems missing	Select **Evidence** for sources, **Review** for notes, or **Output** for shared facts and revision history. Stage selection returns to Output.

Operating limits to state accurately

The app is local and single-user: one API process, one worker, one globally running task, and one active turn per session. Local scope is not multi-user authentication. Do not relax loopback/origin protections for an audience deployment.

Stop is a cooperative cancellation request; wait for its terminal outcome. **Retry** is explicit and creates a linked new task. Failed, cancelled, or stale proposals must not be presented as published output. Closing the browser is not cancellation.

This demo does not detect live anomalies, compute a validated risk score, prove regulatory compliance, or connect to live bank/case systems. Its revision history is not a complete regulated audit-control implementation. The alert list is four packaged fictional alerts: there is no live feed, assignment, triage write, or "no case needed" disposition, and the review-by date is only a reminder. Other scenario features, such as advanced filters and the full operating model, are not implemented.

The matched official A2UI packages retain an upstream Zod dependency advisory, **GHSA-mrhx-c7xc-m3w5**. Fixed-template validation and size/depth limits are safeguards, not a claim that the advisory is resolved. See the README for the dependency disclosure.

Source material and maintenance

README.md is the operational source for installation, configuration, commands, limitations, and validation. The English pair is `demonstration-manual.md` / `demonstration-manual.pdf`; the Korean, Japanese, and Chinese pairs use the `.ko`, `.ja`, and `.zh` suffixes. Each PDF contains its Markdown edition's demonstration content. The translated guides retain English UI labels so presenters can find the actual controls; they do not imply a localized app.

- Scenario narrative and intended process: `docs\scenario\scenario.md` and `docs\scenario\process.md`.
- Implemented design: `docs\superpowers\specs\2026-10-28-investigation-evidence-desk-design.md` and `docs\superpowers\specs\2026-10-29-alert-first-investigation-start-design.md`.
- Protocol design background: `docs\research\ag-ui-a2ui-integration.md`.
- Executable browser acceptance: `web\e2e\investigation-workflow.spec.ts`.

Recheck UI labels and behavior when the app changes. Automated fixture runs verify application behavior, not live inference quality, financial accuracy, or real-system permissions.